

What is R ?

- R is a high-level computer language and environment for statistics and graphics:
 - performs a variety of simple and advanced statistical methods,
 - produces high quality graphics,
 - enables writing new functions that extends R 's uses.
- R is a free open source software maintained by several contributors. Including an “R Core Team” of several programmers who are responsible for modifying the R source code.
- The official R home page is: <http://www.R-project.org>
- The R system can be installed on, Windows, Mac or Linux. Visit the R website and follow the installation directions to install the base system.
- Stock and Watson textbook R companion book:
<https://www.econometrics-with-r.org/index.html>
- Wooldridge textbook R companion book:
<http://www.urfie.net/index.html>

What is R ?

- R is a free implementation of a dialect of the S language. A preliminary version of S was created by Bell Labs in the 1970s that was meant to be a programming language like C but for statistics.
- John Chambers was one of the primary inventors for the language, and he won the Association for Computing Machinery Award in 1999.
- S is a high-level programming language, with similarities to Python. By the 1980s, the popularity of S was growing outside of Bell Labs. This led to Statistical Sciences Inc. creating the S-Plus software package based on S in 1988 and then selling it to users. In 1993, Statistical Sciences Inc. merges with a company named Mathsoft and buys the exclusive license to market S. This prevents anyone from obtaining a free version from Bell Labs. After further company merges, S-Plus now resides with Tibco Software Inc.
- R was initially written by Ross Ihaka and Robert Gentleman at the Department of Statistics of the University of Auckland in Auckland, New Zealand (hence the name).

What is R ?

- Version 1.0.0 of R was released on February 29, 2000. Since this time, the use of R has skyrocketed.
- In addition to the base system there are user contributed add-on packages, i.e., collection of functions, examples and documentation that usually focus on a specific task.
- The base system contains some packages. To install an additional package, say xtable, be connected to the Internet and type

```
install.packages("xtable")
```

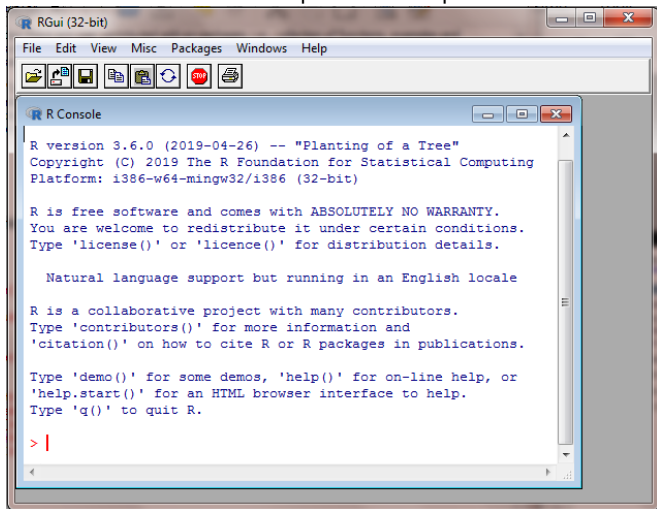
You will be asked to select the mirror site nearest to you, after that everything is automatic.

- In order to use the contents of the package we need to load it,

```
library(xtable) or require(xtable)
```

R Console

- The R Console is where computations are performed.



R as a calculator

- You can use R as a calculator. Enter a math expression and the result is outputted to the console.

Binary Operations	+	-	*	/	^	%%
Math Functions	abs	sqrt	log	exp	log10	factorial
Trig Functions	sin	cos	tan	asin	acos	atan
Rounding	round	ceiling	floor	trunc	signif	zapsmall
Math Quantities	Inf	-Inf	NaN	pi	exp(1)	1i

```
> 5+5*(1+0.4)
[1] 12
> pnorm(1.96)
[1] 0.9750021
> (2-3)/6
[1] -0.1666667
> 2^2
[1] 4
> sin(pi/2)
[1] 1
> log(4)
[1] 1.386294
```

R : Objects and Functions

- What we normally do is create objects and then perform functions on those objects (functions are also considered objects).

- Assign an object a name "x" using either

```
x <- object      (object -> x )
```

```
x = object
```

- Call a function by

function name(list of arguments separated by commas)

- ☐ Each function has a set of formal arguments some with default values; these can be found in the function's documentation.
- ☐ To specify a particular argument use the argument's name.
- R is **CASE SENSITIVE**.

R : Objects and Functions

- Suppose I want to find the mean of a set of numbers. I first assign the vector of numbers to `x` and then call the function `mean()`.

```
> x=c(0,5,7,9,1,2,8)
> x
[1] 0 5 7 9 1 2 8
> mean(x)
[1] 4.571429
> ls()
[1] "x"
> objects()
[1] "x"
```

- Suppose I want to sort a vector `y` so that the numbers are descending.

```
> y=c(4,2,0,9,5,3,10)
> y
[1] 4 2 0 9 5 3 10
> sort(y)
[1] 0 2 3 4 5 9 10
> sort(y,decreasing=TRUE)
[1] 10 9 5 4 3 2 0
```

R : Objects and Functions

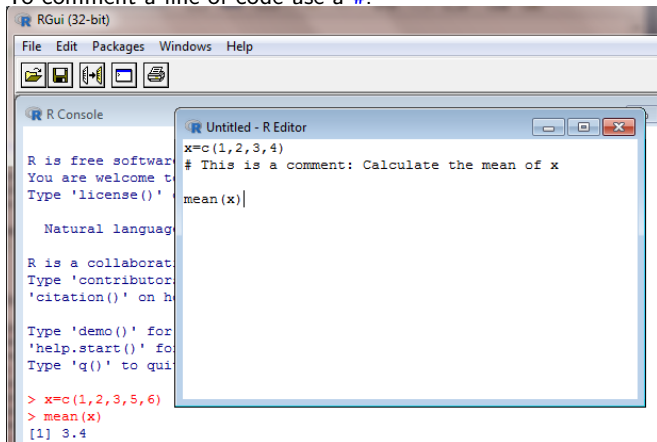
- A function `mysd()` can be written to find the standard deviation of a given vector of numbers.

```
# Declaring a function
mysd<-function(numbers) {
  sqrt(var(numbers))
}

> x<-c(1,2,3,4,5)
# Call the function "mysd()"
> mysd(x)
[1] 1.581139
```


R : Script Editor

- The script editor is used for writing programs in R .
- Starting with R 2.0, a VERY limited program editor was incorporated into R. To open script editor select **File > New script** to create a new program.
- To comment a line of code use a **#**.



R : Script Editor

- Instead of this editor, we will use **RStudio**. This software is available for **free**, and it runs on Linux, Mac, and Windows operating systems.
- This software is actually more than an editor because it integrates a program editor with the **R** Console window, graphics window, R-help window, and other items within one overall window environment. Thus, **RStudio** is an integrated development environment (IDE) for constructing and running **R** programs.
- **RStudio** can be downloaded at <http://www.rstudio.com>.

Documentation

- To look at the documentation for a specific function from a loaded package,
?function name
help(function name)
- For a complete list of functions, use *library(help = package name)*.
- To run the example included with the documentation, *example(function name)*.
- To get an overview of an installed package and its contents,
help(package=package name).
- Some packages include a *vignette*, a file with additional documentation and examples. To get a list of all vignettes from all installed packages, type *vignette()*. To see a particular vignette, *vignette("topic")*.

```
?mean;      ?pnorm;      help("mean");      help("pnorm");  
example("mean");      library(help="stats");  
vignette();      vignette("ggplot2-specs")
```

R : Working Directory & Workspace

- To set the working directory click [Session > Set Working Directory > Choose Directory ...](#) or type `setwd(file path)`.

```
> setwd("C:/Users/????/Desktop/R_intro/R_intro")
> getwd()
[1] "C:/Users/????/Desktop/R_intro/R_intro"
```

- The [workspace](#) is where all the objects you create during a session are located.
- When you close R you will be prompted to “[Save workspace image?](#)” If you say yes, the next time you open R from the same working directory the workspace will be restored.
- That is, all of the objects you created during your last session will still be available.

R : Managing the Workspace

- To list what variables you have created in the current session, `ls()`
- To see what libraries and dataframes are loaded, `search()`
- To see what libraries have been installed, `library()`
- To remove an object, `rm(object names)`
- To remove all objects, `rm(list = ls())`

```
ls() #list of objects created in the current session
search() #loaded libraries and dataframes
library() #installed libraries
```

```
xdata=c(2,30,1,2)
rm(xdata)
```

```
rm(list=ls()) # remove all objects from the current session
```

Vectors

■ Data structures

- ☐ Vectors
- ☐ Matrices
- ☐ Arrays
- ☐ Lists
- ☐ Dataframes

■ Data types

- ☐ Numeric
- ☐ Logical
- ☐ Character
- ☐ Factor
- ☐ Dates
- ☐ Missing data

Vectors

- A vector is an ordered collection of objects of the same type.
- The function `c(...)` concatenates its arguments to form a vector.
- To create a patterned vector
 - `:` Sequence of integers
 - `seq()` General sequence
 - `rep()` Vector of replicated elements

```
v1 <- c(2.5, 4, 7.3, 0.1)      seq(0,2,by=0.5)
v1                               [1] 0.0 0.5 1.0 1.5 2.0
[1] 2.5 4.0 7.3 0.1

                               seq(0,2,len=6)
v2                               [1] 0.0 0.4 0.8 1.2 1.6 2.0
[1] "A" "B" "C" "D"

                               rep(1:5,each=2)
v3                               [1] 1 1 2 2 3 3 4 4 5 5
[1] -3 -2 -1 0 1 2 3

                               rep(1:5,times=2)
                               [1] 1 2 3 4 5 1 2 3 4 5
```

Reference Elements of a Vector

- Use `[]` with a vector/scalar of positions to reference elements of a vector.
- Include a minus sign before the vector/scalar to remove elements.

```
> x <- c(4, 7, 2, 10, 1, 0)
> x[4] # return the fourth element
[1] 10
> x[1:3] # return the first three elements
[1] 4 7 2
> x[c(2,5,6)] # return elements with index c(2,5,6)
[1] 7 1 0
> x[-3] # remove the third element
[1] 4 7 10 1 0
> x[-c(4,5)] # remove the the fourth and fifth element
[1] 4 7 2 0
> x[x>4] # return elements bigger than 4
[1] 7 10
> x[3] <- 99 # change the third element to 99
> x
[1] 4 7 99 10 1 0
```


which() and match()

■ Additional functions that return the indices of a vector:

- ☐ `which()` Indices of a logical vector the condition is **TRUE**
- ☐ `which.max()` Location of the (first) maximum element of a numeric vector
- ☐ `which.min()` Location of the (first) minimum element of a numeric vector
- ☐ `match()` First position of an element in a vector

```
x <- c(4, 7, 2, 10, 1, 0)
x>=4
[1] TRUE TRUE FALSE TRUE FALSE FALSE
which(x>=4)
[1] 1 2 4
which.max(x)
[1] 4
x[which.max(x)]
[1] 10
max(x)
[1] 10
```

```
y <- rep(1:5, times=5:1)
y
[1] 1 1 1 1 1 2 2 2 2 3 3 3 4 4 5
match(1:5, y)
[1] 1 6 10 13 15
match(unique(y), y)
[1] 1 6 10 13 15
```

Useful Vector Functions

<code>sum(x)</code>	<code>prod(x)</code>	Sum/product of the elements of <code>x</code>
<code>cumsum(x)</code>	<code>cumprod(x)</code>	Cumulative sum/product of the elements of <code>x</code>
<code>min(x)</code>	<code>max(x)</code>	Minimum/maximum element of <code>x</code>
<code>mean(x)</code>	<code>median(x)</code>	Mean/median of <code>x</code>
<code>var(x)</code>	<code>sd(x)</code>	Variance/standard deviation of <code>x</code>
<code>cov(x,y)</code>	<code>cor(x,y)</code>	Covariance/correlation of <code>x</code> and <code>y</code>
<code>range(x)</code>		Range of <code>x</code>
<code>quantile(x)</code>		Quantiles of <code>x</code> for the given probabilities
<code>fivenum(x)</code>		Five number summary of <code>x</code>
<hr/>		
<code>length(x)</code>		Number of elements in <code>x</code>
<code>unique(x)</code>		Unique elements of <code>x</code>
<code>rev(x)</code>		Reverse the elements of <code>x</code>
<code>sort(x)</code>		Sort the elements of <code>x</code>
<code>which(x)</code>		Indices of TRUEs in a logical vector
<code>which.max(x)</code>	<code>which.min(x)</code>	Index of the max/min element of <code>x</code>
<code>match(x)</code>		First position of an element in a vector
<hr/>		
<code>union(x,y)</code>		Union of <code>x</code> and <code>y</code>
<code>intersect(x,y)</code>		Intersection of <code>x</code> and <code>y</code>
<code>setdiff(x,y)</code>		Elements of <code>x</code> that are not in <code>y</code>
<code>setequal(x,y)</code>		Do <code>x</code> and <code>y</code> contain the same elements?

Matrices

- To create a matrix

```
matrix(data=NA, nrow=1, ncol=1, byrow = FALSE, dimnames = NULL)
```

data a vector that gives data to fill the matrix; if data does not have enough elements to fill the matrix, then the elements are recycled

nrow desired number of rows

ncol desired number of columns

byrow if **FALSE** (default) matrix is filled by columns, otherwise by rows

dimnames (optional) list of length 2 giving the row and column names respectively, list names will be used as names for the dimensions

```
x=matrix(c(5,0,6,1,3,5,9,5,7,1,5,3), nrow=3, ncol=4, byrow=TRUE,  
          dimnames=list(rows=c("r.1", "r.2", "r.3"),  
                          cols=c("c.1", "c.2", "c.3", "c.4")))
```

```
x  
      cols  
rows  c.1 c.2 c.3 c.4  
r.1    5  0  6  1  
r.2    3  5  9  5  
r.3    7  1  5  3
```

Reference Elements of a Matrix

- Reference matrix elements using the `[]` just like with vectors, but now with 2-dimensions

```
x=matrix(c(5,0,6,1,3,5,9,5,7,1,5,3), nrow=3, ncol=4, byrow=TRUE)
x
```

	[,1]	[,2]	[,3]	[,4]
[1,]	5	0	6	1
[2,]	3	5	9	5
[3,]	7	1	5	3

```
x[2,3] # Row 2, Column 3
```

```
[1] 9
```

```
x[1,] # Row 1
```

```
[1] 5 0 6 1
```

```
x[,2] # Column 2
```

```
[1] 0 5 1
```

```
x[c(1,3),] # Rows 1 and 3, all Columns
```

	[,1]	[,2]	[,3]	[,4]
[1,]	5	0	6	1
[2,]	7	1	5	3

Matrix operations

- When matrices are used in math expressions the operations are performed element by element.
- For matrix multiplication use the `%*%` operator
- Using `%*%` on two vectors will return the inner product as a matrix and not a scalar. Use `c()` or `as.vector()` to convert to a scalar.

```
A<-matrix(1:4, nrow=2)
B<-matrix(1, nrow=2, ncol=2)
A*B
A%*%B
      [,1] [,2]
[1,]     4     4
[2,]     6     6
y <- 1:3
y%*%y # returns dot product in matrix form
      [,1]
[1,]    14
A/(y%*%y)
Error in A/(y %*% y) : non-conformable arrays
A/c(y%*%y)
      [,1] [,2]
[1,] 0.07142857 0.2142857
[2,] 0.14285714 0.2857143
```

Useful Matrix Functions

<code>t(A)</code>	Transpose of A
<code>det(A)</code>	Determinant of A
<code>solve(A,b)</code>	Solves the equation $Ax=b$ for x
<code>solve(A)</code>	Matrix inverse of A
<code>MASS::ginv(A)</code>	Generalized inverse of A (MASS Package)
<code>eigen(A)</code>	Eigenvalues and eigenvectors of A
<code>chol(A)</code>	Cholesky factorization of A
<code>diag(n)</code>	Creates $n \times n$ identity matrix
<code>diag(A)</code>	Returns the diagonal elements of a matrix A
<code>diag(x)</code>	Create a diagonal matrix from a vector x
<code>lower.tri(A), upper.tri(A)</code>	Matrix of logicals indicating lower/upper triangular matrix
<hr/>	
<code>apply</code>	Apply a function to the margins of a matrix
<code>rbind(...)</code>	Combines arguments by rows
<code>cbind(...)</code>	Combines arguments by columns
<code>dim(A)</code>	Dimensions of A
<code>nrow(A), ncol(A)</code>	Number of rows/columns of A
<code>colnames(A), rownames(A)</code>	Get or set the column/row names of A
<code>dimnames(A)</code>	Get or set the dimension names of A

Arrays

- An array is a multi-dimensional generalization of a vector
- To create an array

```
array(data = NA, dim = length(data), dimnames = NULL)
```

data A vector that gives data to fill the array; if data does not have enough elements to fill the array, then the elements are recycled

dim Dimension of the array, a vector of length one or more giving the maximum indices in each dimension

dimnames Name of the dimensions, list with one component for each dimension, either NULL or a character vector of the length given by dim for that dimension. The list can be named, and the list names will be used as names for the dimensions.

- Values are entered by columns
- Like with vectors and matrices, when arrays are used in math expressions the operations are performed element by element.
- Also like vectors and matrices, the elements of an array must all be of the same type (numeric, character, logical, etc.)

Arrays

- Sample $2 \times 3 \times 2$ array

```
w <- array(1:12, dim=c(2,3,2)  
          dimnames=list(c("A","B"), c("X","Y","Z"), c("N","M")))
```

```
w  
 , , N
```

	X	Y	Z
A	1	3	5
B	2	4	6

```
 , , M
```

	X	Y	Z
A	7	9	11
B	8	10	12

- You can think of "N", "M" as page numbers. Elements of an array can be referenced using `[]`.

Arrays

```
w[2,3,1] # Row 2, Column 3, Matrix 1
[1] 6
```

```
w[, "Y", ] # Column named "Y"
      N    M
A 3    9
B 4   10
```

```
w[1,,] # Row 1
      N    M
X 1    7
Y 3    9
Z 5   11
```

```
w[1:2,,"M"] # Rows 1 and 2, Matrix "M"
      X    Y    Z
A 7    9   11
B 8   10   12
```

Lists

- A list is a general form of a vector, where the elements don't need to be of the same type or dimension.
- The function `list(...)` creates a list of the arguments.
- Arguments have the form *name=value*. Arguments can be specified with and without names.

```
x <- list(num=c(1,2,3), "Nick", identity=diag(2))
```

```
x
```

```
$num
```

```
[1] 1 2 3
```

```
[[2]]
```

```
[1] "Nick"
```

```
$identity
```

	[,1]	[,2]
[1,]	1	0
[2,]	0	1

- Elements of a list can be referenced using `[]` as well as `[[]]` or `$`.

Lists

```
y=list(v1=c(2,3,5),"Econometrics",A=matrix(1:4, nrow=2))
```

```
y[[2]] # Second element of y  
[1] "Econometrics"
```

```
y[["v1"]] # Element named "v1"  
[1] 2 3 5
```

```
y$A # Element named "A"  
      [,1] [,2]  
[1,]    1    3  
[2,]    2    4
```

```
y[[3]][1,] # First row of the third element  
[1] 1 3
```

```
y[1:2] # Create a sub-list of the first two elements  
$v1  
[1] 2 3 5
```

```
[[2]]  
[1] "Econometrics"
```

Logical

- Logical values are represented by the reserved words **TRUE** and **FALSE** in all caps or simply **T** and **F**.

<code>!x</code>	NOT <code>x</code>
<code>x & y</code>	<code>x</code> and <code>y</code> elementwise, returns a vector
<code>x && y</code>	<code>x</code> and <code>y</code> returns a single value
<code>x y</code>	<code>x</code> OR <code>y</code> elementwise, returns a vector
<code>x y</code>	<code>x</code> OR <code>y</code> returns a single value
<code>xor(x,y)</code>	Exclusive OR of <code>x</code> and <code>y</code> , elementwise
<code>x %in% y</code>	<code>x</code> IN <code>y</code>

<code>x < y</code>	<code>x < y</code>
<code>x > y</code>	<code>x > y</code>
<code>x <= y</code>	<code>x ≤ y</code>
<code>x >= y</code>	<code>x ≥ y</code>
<code>x == y</code>	<code>x = y</code>
<code>x != y</code>	<code>x ≠ y</code>

<code>isTRUE(x)</code>	TRUE if <code>x</code> is TRUE
<code>all(...)</code>	TRUE if all arguments are TRUE
<code>any(...)</code>	TRUE if at least one argument is TRUE
<code>identical(x,y)</code>	Safe and reliable way to test two objects for being EXACTLY equal
<code>all.equal(x,y)</code>	Test if two objects are NEARLY equal

Logical

```
x <- 1:10
(x%2==0) | (x>5) # What elements of x are even or greater than 5
[1] FALSE TRUE FALSE TRUE FALSE TRUE TRUE TRUE TRUE

y <- 5:15 # What elements of x are in y
x %in% y
[1] FALSE FALSE FALSE FALSE TRUE TRUE TRUE TRUE TRUE

x[x %in% y]
[1] 5 6 7 8 9 10

any(x>5) # Are any elements of x greater than 5?
[1] TRUE

all(x>5) # Are all elements of x greater than 5?
[1] FALSE
```

Isn't that equal?

- In general, logical operators may not produce a single value and may return an **NA** if an element is **NA** or **NaN**.
- If you must get a single **TRUE** or **FALSE**, such as with if expressions, you should NOT use **==** or **!=**. Unless you are absolutely sure that nothing unusual can happen, you should use the **identical()** function instead.
- **identical()** only returns a single logical value, **TRUE** or **FALSE**, never **NA**.

```
name <- "Nick"
if(name=="Nick") TRUE else FALSE
[1] TRUE
```

- What if name is never set to "Nick"?

```
name <- NA
if(name=="Nick") TRUE else FALSE
Error in if (name=="Nick") TRUE else FALSE :
missing value where TRUE/FALSE needed
```

```
if(identical(name, "Nick")) TRUE else FALSE
[1] FALSE
```

Isn't that equal?

- With `all.equal()` objects are treated as equal if the only difference is probably the result of inexact floating-point calculations. Returns `TRUE` if the mean relative difference is less than the specified tolerance.
- `all.equal()` either returns `TRUE` or a character string that describes the difference. Therefore, do not use `all.equal()` directly in if expressions, instead use with `isTRUE()` or `identical()`.

```
(x <- sqrt(2))  
[1] 1.414214  
x^2  
[1] 2  
x^2==2  
[1] FALSE  
all.equal(x^2,2)  
[1] TRUE  
all.equal(x^2,1)  
[1] "Mean relative difference: 0.5"  
isTRUE(all.equal(x^2,1))  
[1] FALSE
```

Character

- Character strings are denoted by quotation marks, single ' ' or double " "

<code>cat()</code>	Concatenate objects and print to console(\n new line)
<code>paste()</code>	Concatenate objects and return a string
<code>print()</code>	Print an object
<code>substr()</code>	Extract or replace substrings in a character vector
<code>strtrim()</code>	Trim character vectors to specified display widths
<code>strsplit</code>	Split elements of a character vector according to a substring
<code>grep()</code>	Search for matches to a pattern within a character vector
<code>grepl()</code>	Like <code>grep()</code> , but returns a logical vector
<code>agrep()</code>	Similar to <code>grep()</code> , but searches for approximate matches
<code>regexpr()</code>	Similar to <code>grep()</code> , but returns the position of the first instance of a pattern within a string
<code>gsub()</code>	Replace all occurrences of a pattern with a character vector
<code>sub()</code>	Like <code>gsub()</code> , but only replaces the first occurrence
<code>tolower(), toupper()</code>	Convert to all lower/upper case
<code>noquote()</code>	Print a character vector without quotations
<code>nchar()</code>	Number of characters
<code>letters, LETTERS</code>	Built-in vector of lower and upper case letters

Character

```
animals=c("bird", "horse", "fish")
home=c("tree", "barn", "lake")

length(animals) # Number of strings
nchar(animals) # Number of characters in each string

cat("Animals:", animals) # Need \n to move cursor to a newline
cat(animals, home, "\n") # Joins one vector after the other

paste(animals, collapse=" ")#Create one long string of animals
a.h=paste(animals,home,sep=".")#Pairwise joining of animals and home

# Split strings at ".", fixed=TRUE since "." is used
# for pattern matching
unlist(strsplit(a.h, ".", fixed=TRUE))
substr(animals, 2, 4) # Get characters 2-4 of each animal
strtrim(animals, 3) # Print the first three characters
toupper(animals) # Print animals in all upper case
```

Factor

- A factor is a categorical variable with a defined number of ordered or unordered levels.

```
x=factor(rep(1:3, 4), labels=c("low", "med", "high"), ordered=TRUE)
x
[1] low  med  high low  med  high low  med  high low  med  high
Levels: low < med < high
```

<code>levels(x)</code>	Retrieve or set the levels of x
<code>nlevels(x)</code>	Return the number of levels of x
<code>relevel(x, ref)</code>	Levels of x are reordered so that the level specified by ref is first
<code>reorder()</code>	Reorders levels based on the values of a second variable
<code>gl()</code>	Generate factors by specifying the pattern of their levels
<code>cut(x, breaks)</code>	Divides the range of x into intervals (factors) determined by breaks

```
> nlevels(x)
[1] 3
> levels(x)
[1] "low" "med" "high"
> is.factor(x)
[1] TRUE
> is.ordered(x)
[1] TRUE
```

Statistical Distributions

- R has several built-in statistical distributions and for each distribution four functions are available:
 - ☐ **r** Random number generator
 - ☐ **d** Density function
 - ☐ **p** Cumulative distribution function
 - ☐ **q** Quantile function
- Each letter can be added as a prefix to any of the **R** distribution names.

R	Distribution	R	Distribution
beta	Beta	logis	Logistic
binom	Binomial	nbinom	Negative binomial
cauchy	Cauchy	norm	Normal
exp	Exponential	pois	Poisson
chisq	Chi-square	signrank	Wilcoxon signed rank statistic
f	Fisher's F	t	Student's t
gamma	Gamma	unif	Uniform
geom	Geometric	weibull	Weibull
hyper	Hypergeometric	wilcox	Wilcoxon rank sum
lnorm	Lognormal		

Statistical distributions

```
> dnorm(1.96, mean=0, sd=1) # Density
[1] 0.05844094
> pnorm(1.96, mean=0, sd=1) # Distribution (lower tail)
[1] 0.9750021
# Distribution(upper tail)
> pnorm(1.96, mean=0, sd=1, lower.tail=FALSE)
[1] 0.0249979
> qnorm(0.975, mean=0, sd=1) # Quantile
[1] 1.959964
> rnorm(5, mean=0, sd=1) # Five random numbers from N(0,1)
[1] -0.1921395 1.0880367 1.2215088 1.2261773 -0.2187957

> set.seed(17632)
> runif(5)
[1] 0.03288684 0.88861821 0.21466744 0.32907126 0.78222193
> set.seed(17632)
> runif(5)
[1] 0.03288684 0.88861821 0.21466744 0.32907126 0.78222193
```

Control Structures

- See `?Control` for the R documentation on control structures.
 - ☐ `if()...else` Determine which set of expressions to run depending on whether or not a condition is TRUE
 - ☐ `ifelse()` Do something to each element in a data structure depending on whether or not a condition is TRUE for that particular element
 - ☐ `switch()` Evaluate different expressions depending on a given value
 - ☐ `for()` Loop for a fixed number of iterations
 - ☐ `while()` Loop until a condition is FALSE
 - ☐ `repeat` Repeat until iterations are halted with a call to break
 - ☐ `break` Break out of a loop
 - ☐ `next` Stop processing the current iteration and advance the looping index

if()...else

- The condition needs to evaluate to a single logical value.
- Brackets `{}` are not necessary if you only have one expression and/or the `if()...else` statement are on one line.
- To avoid a syntax error, you should NOT have a newline between the closing bracket of the `if` statement and the `else` statement.

`if(condition) expression if TRUE`

```
if(condition) {  
    expression if TRUE  
}
```

`if(condition) expression if TRUE else expression if FALSE`

```
if(condition) {  
    expression if TRUE  
}else{  
    expression if FALSE  
}
```

Example if()...else

- Calculate the median of a random sample X_1, \dots, X_n

$$\text{median}(X) = \begin{cases} \frac{1}{2}X_{(\frac{n}{2})} + \frac{1}{2}X_{(1+\frac{n}{2})} & \text{if } n \text{ is even} \\ X_{(\frac{n+1}{2})} & \text{if } n \text{ is odd} \end{cases}$$

where $X_{(1)}, \dots, X_{(n)}$ denote order statistics.

```
x <- c(5,4,2,8,9,10)
```

```
n <- length(x)
```

```
sort.x <- sort(x)
```

```
if(n%%2==0){  
  median <- (sort.x[n/2]+sort.x[1+n/2])/2  
} else {  
  median <- sort.x[(n+1)/2]  
}
```

```
median
```

Example: ifelse()

- `ifelse()` returns a value with the same structure as `test` which is filled with elements selected from either `yes` or `no` depending on whether the element of `test` is **TRUE** or **FALSE**.

```
(x <- seq(0,2,len=6))
```

```
[1] 0.0 0.4 0.8 1.2 1.6 2.0
```

```
ifelse(x <= 1, "small", "big")
```

```
[1] "small" "small" "small" "big" "big" "big"
```

```
(y <- matrix(1:8, nrow=2))
```

	[,1]	[,2]	[,3]	[,4]
[1,]	1	3	5	7
[2,]	2	4	6	8

```
ifelse(y>3 & y<7, 1, 0)
```

	[,1]	[,2]	[,3]	[,4]
[1,]	0	0	1	0
[2,]	0	1	1	0

Example: switch()

- The switch() function evaluates and returns different expressions depending on the value of expr: `switch(expr, ...)`.
- expr needs to be a single character string or a number, not a vector.
- ... is a comma separated list of name=expression or number=expression. If there is no expression, then the next non-missing expression is used.

```
central <- function(y, measure) {  
  switch(measure,  
    Mean = ,  
    mean = mean(y),  
    median = median(y),  
    geometric = prod(y)^(1/length(y)),  
    "Invalid Measure")  
}  
  
> y=runif(100)  
> central(y, "mean")  
[1] 0.4844221  
> central(y, "Mean")  
[1] 0.4844221  
> central(y, "Median")  
[1] "Invalid Measure"  
> central(y, "geometric")  
[1] 0.3490125
```

Example: for loops

```
for (var in seq) {  
  expressions  
}
```

- var is the name of the loop variable that changes with each iteration.
- With each iteration var takes on the next value in seq. At the end of the loop var will equal the last value of seq. The number of iterations equals the length of seq.
- Simulating AR(2): $y_t = \alpha_1 y_{t-1} + \alpha_2 y_{t-2} + u_t$, where $\alpha_1 = 0.9$, $\alpha_2 = -0.2$ and $u_t \sim N(0, 1)$.

```
# Simulating AR(2) model  
n=1000  
u=rnorm(n)  
y=rep(0,n)  
for(i in 3:1000){  
  y[i]=0.9*y[i-1]-0.2*y[i-2]+u[i]  
}  
y=y[501:1000] #take the last 500 observation  
plot(y,type="l")
```

Example: while loops

```
while (cond) {
  expressions
}
```

- cond is a single logical value that is not NA.

```
# Calculate 10! using a while loop
```

```
i <- 10
```

```
f <- 1
```

```
while(i>1) {
```

```
  f <- i*f
```

```
  i <- i-1
```

```
  cat(i, f, "\n")
```

```
}
```

```
> f
```

```
[1] 3628800
```

```
> factorial(10)
```

```
[1] 3628800
```

Example: repeat loops

```
repeat {  
  expressions  
  if(cond) break  
}
```

- The repeat loop does not contain a limit. Therefore it is necessary to include an if statement with the break command to make sure you do not have an infinite loop.

```
# Calculate 10! using a repeat loop
```

```
i <- 10  
f <- 1  
repeat{  
  f <- i*f  
  i <- i-1  
  cat(i, f, "\n")  
  if(i<1){  
    break  
  }  
}
```

```
> f  
[1] 3628800  
> factorial(10)  
[1] 3628800
```

Vectorizing computations

- Functions that perform whole object computations or vectorizing computations can improve efficiency.

<code>lapply()</code> , <code>sapply()</code>	Apply a function to each element of a vector or list
<code>mapply()</code>	Multivariate version of <code>sapply()</code> , apply a function to the corresponding elements from multiple vectors
<code>Vectorize()</code>	Return a new function that acts as if <code>mapply()</code> was called

- These approaches may not improve efficiency because they do not reduce the number of function calls.

sapply()

```
poisson <- function(x, lambda){
  exp(-lambda)*lambda^x/factorial(x)
}
poisson(1:5, lambda=1)
# Suppose lambda=1,2,3,4,5. We can use a loop
result=matrix(nrow=5, ncol=5,
dimnames = list(rows=c("x=1","x=2","x=3","x=4","x=5"),
cols=c("lam=1","lam=2","lam=3","lam=4","lam=5")))
for(lambda in seq(1:5)){
  result[,lambda]=poisson(1:5,lambda)
}
result
  cols
rows   lam=1   lam=2   lam=3   lam=4   lam=5
x=1 0.367879441 0.27067057 0.1493612 0.07326256 0.03368973
x=2 0.183939721 0.27067057 0.2240418 0.14652511 0.08422434
x=3 0.061313240 0.18044704 0.2240418 0.19536681 0.14037390
x=4 0.015328310 0.09022352 0.1680314 0.19536681 0.17546737
x=5 0.003065662 0.03608941 0.1008188 0.15629345 0.17546737

# Alternatively, we can use sapply()
sapply(c(lam1=1,lam2=2,lam3=3,lam4=4,lam5=5), poisson, x=1:5)
```

mapply()

```
# Evaluating gamma density
gamma.density <- function(x, alpha, beta) {
  x^(alpha-1)*exp(-x/beta)/gamma(alpha)/beta^alpha
}

# Evaluate the gamma density at 1:5 for
# alpha=beta=1, alpha=beta=2, alpha=beta=3
mapply(gamma.density, alpha=1:3, beta=1:3, MoreArgs=list(x=1:5))

# Evaluate the gamma density at 1:5 for
# all combinations of alpha and beta
parm <- expand.grid(alpha=1:3, beta=1:3)
out <- mapply(gamma.density, parm[,1], parm[,2],
              MoreArgs=list(x=1:5))
colnames(out)=paste("alpha",parm[,1],"beta",parm[,2],sep=".")
out
```

apply()

```
> # apply()
> A <- matrix(sample(1:100, 12), nrow=3, ncol=4)

> # Calculate the mean of each column
> apply(A, 2, FUN=mean)
[1] 43.66667 58.00000 52.33333 52.33333

> # Simply functions can be defined within the apply() command
> # Calculate variance of each column
> apply(A, 2, FUN=function(x) sum((x-mean(x))^2)/(length(x)-1))
[1] 2012.3333 1764.0000 192.3333 569.3333
```


Functions

- To declare a function:

```
function name <- function(argument list) {  
  body  
}
```

- The argument list is a comma-separated list of formal argument names: (i) name, (ii) name=default value, (iii)
- The ... is a list of the remaining arguments that do not match any of the formal arguments.
- Generally, the body is a group of R expressions contained in curly brackets {}. If the body is only one expression the curly brackets are optional.

```
myfunction<-function(numbers) {  
  x1=sd(numbers) # calculate sd  
  x2=mean(numbers) # calculate mean  
  x3=median(numbers) # calculate mean  
  out=list(mean=x2,sd=x1,median=x3)  
  return(out)  
}  
x=c(3,4,5,6,7,8,1)  
myfunction(x)
```

Functions

- Example: Consider a random sample x_1, x_2, \dots, x_n with a sample mean \bar{x} and sample variance s . The $100(1 - \alpha)\%$ t-confidence interval for the mean is

$$\bar{x} \pm t_{1-\alpha/2} \times \frac{s}{\sqrt{n}},$$

where $t_{1-\alpha/2}$ is the critical value for the t distribution with $(n - 1)$ degrees of freedom.

- If x_1, x_2, \dots, x_n is a random sample from a normal population then the t-confidence interval of the mean is exact and is approximately correct for large n otherwise.
- So how large of an n is necessary?
- To investigate this question, we will write a R function that performs a simulation for a given statistical distribution and a sample size.

Functions

```
sim.t.CI <- function(RVgen,n,nsim=1000,alpha=.05,mu,...){  
  # Inputs:  
  # RVgen: Random variable generator, i.e. rnorm  
  # n: Sample size, possibly a vector  
  # nsim=1000, number of iterations, default 1000  
  # alpha=.05, 100(1-alpha) confidence level  
  # mu: Population mean  
  # ... Arguments to be passed to RVgen  
  CP <- array(dim=nsim)  
  for(j in 1:nsim) {  
    # Generate random sample  
    x <- RVgen(n, ...)  
    # Calculate t-confidence interval  
    xbar <- mean(x)  
    lower <- xbar- qt(1-alpha/2, n-1)*sqrt(var(x)/n)  
    upper <- xbar+ qt(1-alpha/2, n-1)*sqrt(var(x)/n)  
    CP[j] <- ifelse(lower < mu & mu < upper, 1, 0)  
  }  
  result=mean(CP)  
  return(result)  
}
```

Functions

```
> sim.t.CI(rnorm, n=30, mu=0, mean=0, sd=1)
[1] 0.952
> sim.t.CI(rnorm, n=100, mu=0, mean=0, sd=1)
[1] 0.952

> sim.t.CI(rchisq, n=30, mu=3, df=3)
[1] 0.936
> sim.t.CI(rchisq, n=100, mu=3, df=3)
[1] 0.945

> sim.t.CI(runif, n=5, mu=.5, min=0, max=1)
[1] 0.933
> sim.t.CI(runif, n=100, mu=.5, min=0, max=1)
[1] 0.958
```